

Anthony Poerio (adp59@pitt.edu)

CS1674: Homework 11 - Written

Due: 12/7/2016, 11:59pm

PART I

1) Compute network activations using fixed input and weights.

First, we need to compute Z_2 . This can be done with the following formula (after expanding the summation and removing the bias value):

$$\mathbf{a}_2 = (w_{21}^{(1)} * x_1) + (w_{22}^{(1)} * x_2) + (w_{23}^{(1)} * x_3) + (w_{24}^{(1)} * x_4)$$

$$\mathbf{a}_2 = (0.02 * x_1) + (0.25 * x_2) + (0.4 * x_3) + (0.3 * x_4)$$

$$\mathbf{a}_2 = (0.02 * 10) + (0.25 * 1) + (0.4 * 2) + (0.3 * 3)$$

$$\mathbf{a}_2 = 2.15$$

Next, perform tanh activation.

$$\mathbf{Z}_2 = \tanh(\mathbf{a}_2)$$

$$\mathbf{Z}_2 = \tanh(2.15)$$

$$\mathbf{Z}_2 = \mathbf{0.9732}$$

PART II

What is the output size resulting from convolving a 35x35 image with a filter of 15x15, using:

a) Stride 1 and no padding

- i) Formula to use: $(N-F) / \text{stride} + 1$
 - o $N=35$

- $F=15$
- $\text{Stride}=1$
- $\text{Padding}=0$
- $\text{Formula} = (35-15) / 1 + 1$
 - $20/1 + 1 = 21$
 - Therefore, **filter size = 21x21**

b) Stride 1 and padding 1

- i) Formula to use: $(N-F) / \text{stride} + 1$
 - $N=35$
 - $F=15$
 - $\text{Stride}=1$
 - $\text{Padding} = 2 \text{ dimensions} * 1 \text{ px} = (2*1)$
 - $\text{Formula} = ((35+(2*1))-15) / 1 + 1$
 - $22/1 + 1 = 23$
 - Therefore, **filter size = 23x23**

c) Stride 2 and padding 3

- i) Formula to use: $(N-F) / \text{stride} + 1$
 - $N=35$
 - $F=15$
 - $\text{Stride}=1$
 - $\text{Padding} = 2 \text{ dimensions} * 3 \text{ px} = (2*3)$
 - $\text{Formula} = ((35+(2*3))-15) / 1 + 1$
 - $(41-15)/1 + 1$
 - $26/1 + 1 = 27$
 - Therefore, **filter size = 27x27**

PART III

Computing outputs of convolutions

a) First, show the output of applying convolution

To start, let's figure out the dimensions of our resulting matrix. The formula to use is:
 $(N-F) / \text{Stride} + 1$

And we know this about the data set:

- Padding=0
- Stride=0
- N=9
- F=3
 - Therefore:
 - $= (N-F) / \text{Stride} + 1$
 - $= (9-3)/2 + 1$
 - $= 6/2 + 1$
 - $= 3 + 1$
 - $= 4$
 - **= 4x4 matrix**

With this in mind, we can calculate the values for each cell of our matrix, given our **Image** and **Filter**, as specified by the assignment prompt.

Calculating by hand, I got:

$[-2 \ -2 \ -1 \ 0; 0 \ 0 \ -3 \ -1; 0 \ 0 \ 0 \ -3; 0 \ 0 \ 0 \ 0]$

Note, that I am doing a convolution by hand, using the formula on slide 26, here:

https://people.cs.pitt.edu/~kovashka/cs1674/vision_04_filters_texture.pdf

b) Second, show the output of applying a Rectified Linear Unit (ReLU) activation

Next, we apply a ReLU activation, which yields a matrix of all 0's. More specifically, our matrix becomes: $[0 \ 0 \ 0 \ 0; 0 \ 0 \ 0 \ 0; 0 \ 0 \ 0 \ 0; 0 \ 0 \ 0 \ 0]$

This is because ReLU transforms any value that is ≤ 0 to 0.

And, all of our values in the convolution matrix produced in part (a) are ≤ 0 .

c) Third, show the output of applying max pooling over 2x2 regions

This yields a matrix with almost all 1's and only one 0, in the top rightmost cell:

[1 1 1 0; 1 1 1 1; 1 1 1 1; 1 1 1 1]

This is because max-pooling takes the MAX value from each region. And the only values in our matrix are 0's and 1's. AND \rightarrow the top rightmost 2x2 region is the only one that does not have a 1 contained in the 2x2 filter.

PART IV

Compute 2 types of loss functions: SVM and Softmax. Use three sets of weights W, each will result in a different set of scores, S \rightarrow for 4 image examples. Goal is: determine which set of weights results in the smallest SVM loss, and which set results in the smallest Softmax loss.

i) **SVM_LOSS**, Results are:

- W1_Loss=6.5676
- **W2_Loss=4.8363 \rightarrow W2 is has the SMALLEST loss.**
- W3_Loss=7.7445

ii) **SOFTMAX_LOSS**, Results are (all ended up negative):

- **W1_Loss= -22.1060 \rightarrow W1 is has the SMALLEST loss.**
- W2_Loss= -7.6050
- W3_Loss= -19.1920

PART V

Compute the numerical gradient of W1.

h=0.0001 yields this gradient vector for me:

G =

1.0e+03 *

-7.0926

-7.5459

-2.7593

-6.7960

-6.5500

-1.6251

-1.1890

-4.9826

-9.5964

-3.4029

-5.8517

-2.2371

-7.5117

-2.5500

-5.0586

-6.9898

-8.9080

-9.5919

-5.4712

-1.3852

-1.4919

-2.5741

-8.4062

-2.5418

-8.1418

-2.4342

-9.2916

-3.4988

-1.9650

-2.5098

-6.1594

-4.7319

-3.5156

-8.3073

-5.8516

-5.4962

-9.1709

-2.8574

-7.5710

-7.5363

-3.8035

-5.6772

-0.7575

-0.5385

-5.3070

-7.7907

-9.3391

-1.2981

-5.6872

-4.6929

-0.1180

-3.3702

-1.6208

-7.9418

-3.1112

-5.2843

-1.6555

-6.0188

-2.6287

-6.5398

-6.8911

-7.4805

-4.5044

-0.8372

-2.2888

-9.1324

-1.5228

-8.2572

-5.3824

-9.9603

-0.7808

-4.4258

-1.0655

-9.6180

-0.0453

-7.7481

-8.1720

-8.6859

-0.8434

-3.9968

-2.5977

-7.9997

-4.3131

-9.1055

-1.8175

-2.6370

-1.4544

-1.3597

-8.6919

-5.7960

-5.4976

-1.4485

-8.5293

-6.2196

-3.5085

-5.1315

-4.0171

-0.7587

-2.3982

-1.2322

And computing the SVM Loss on each example yields these loss values for me:

- $\text{Loss_X1} = 1.2224\text{e}+03$
- $\text{Loss_X2} = 0$
- $\text{Loss_X3} = 1.4240\text{e}+04$
- $\text{Loss_X4} = 8.8725\text{e}+03$

Then, after updating the data with $h=0.001$, I see the following results.

- **Loss_X1 = 1.3312**
- **Loss_X2 = 0.6210**
- **Loss_X3 = 0.0966**
- **Loss_X4 = 1.0404**