

The background image is a vibrant coastal scene. In the upper left, a cluster of buildings with red-tiled roofs sits atop a steep, rocky cliff. The cliff face is light-colored and shows some greenery. Below the cliff, the sea is a deep blue-green, with white foam from waves crashing against dark, moss-covered rocks in the foreground. The sky is a bright blue, filled with fluffy white clouds. A small sailboat is visible on the horizon to the right.

# *CS 1674: Intro to Computer Vision*

# **Matlab Tutorial**

Prof. Adriana Kovashka  
University of Pittsburgh  
August 31, 2016

# Plan for Today

- Intro
  - Course basics refresher
  - Overview of topics
  - What are images?
  - Linear algebra lightning-quick over-/re-view
- **Matlab tutorial**
- Outro
  - Overview of Homework 1 W and P (if time)

# Course Info

- **Course website:**

<http://people.cs.pitt.edu/~kovashka/cs1674>

- **Instructor:** Adriana Kovashka

(kovashka@cs.pitt.edu)

– Please use "CS1674" at the beginning of your Subject

- **Office:** Sennott Square 5129

- **Office hours:** MW, 3:30pm - 4:25pm



# TAs

- **TA/Grader:** Yuhuan Jiang  
(yuhuan@cs.pitt.edu)
- **Additional TAs (office hours only):** Chris Thomas and Nils Murrugarra
- **TAs' office hours:** TBD
- **Your homework:** Fill out the Doodle at <http://doodle.com/poll/gskprtb5uq5k85bm>  
(ignore dates, look at days of the week, 30-min increments)

# Matlab

- You can get it for free through [my.pitt.edu](http://my.pitt.edu) → My Resources → Software Downloads
- Get the latest version (mostly because that's what I use)
- Make sure to check the “Image Processing Toolbox”, “Computer Vision System Toolbox”, and “Statistics and Machine Learning Toolbox” boxes during installation; easiest to install all

# Course Components

- Written HW (11 assignments x 1% each = 11%)
- Programming HW (11 assignments x 4% each = 44%)
- Midterm exam (15%)
- Final exam (25%)
- Participation (5%)

# The Rest of the Course Policies...

- Read the course website carefully!

# Warnings



# Warning #1

- This class is **a lot of work**
- This time I've opted for shorter, more manageable HW assignments, but there is more of them
- I expect you'd be spending **6-8 hours** on homework each week
- ... But you get to understand algorithms and concepts in detail!

# Warning #2

- Some parts will be **hard** and require that you pay close attention!
- ... I will use the written HW to gauge how you're doing
- ... I will also pick on students randomly to answer questions
- **Use instructor's and TAs' office hours!!!**
- ... You will learn a lot!

# Warning #3

- Programming assignments will be in Matlab since that's very common in computer vision, and is optimized for work with matrices
- Matlab also has great documentation
- HW1P is just Matlab practice
- Some people **won't like Matlab** (I like it!)
- ... You will learn a new programming language!

# Clarification from last time

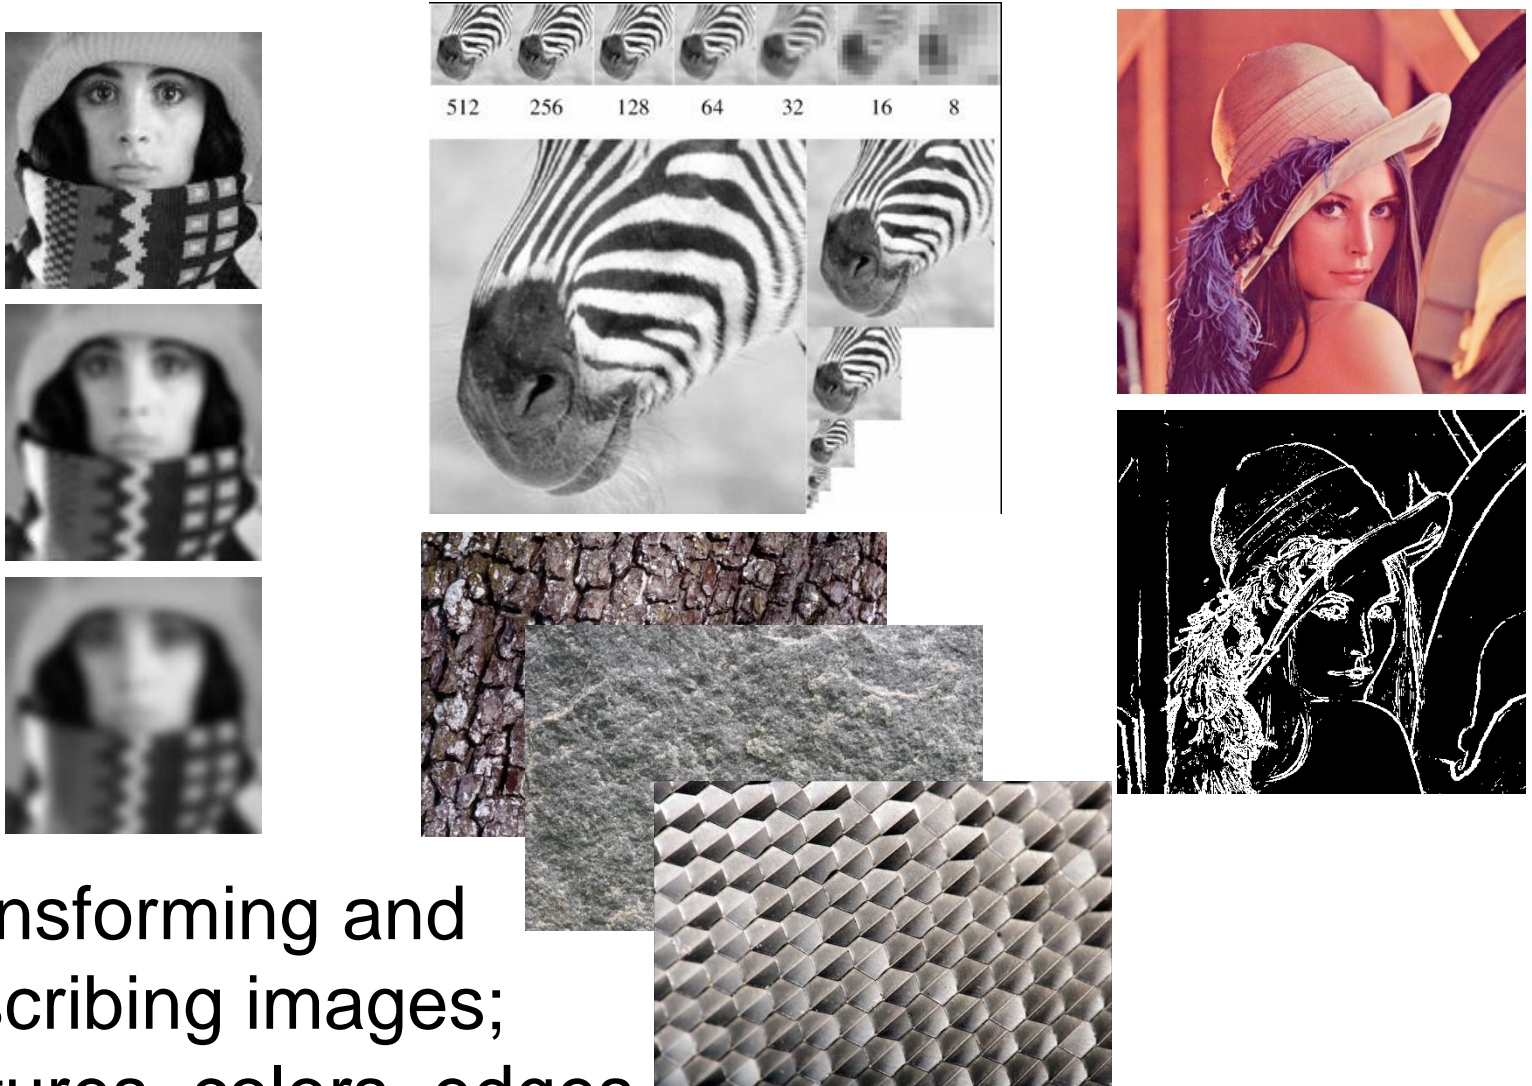
- What is the role of datasets?
  - Whatever our computer vision algorithms learn, they will learn from some (set of) datasets of images; we will also use the datasets to test our algorithms
- Why did I say it's hard to understand deep learning (a.k.a. deep neural networks)?
  - Because most deep learning methods just take an image as input and output predictions, and they learn how to represent and examine the image *on their own*, so they appear to be “black boxes”

Questions?

# Overview of Topics

The next 13 slides will be very quick, then we'll slow down.  
Ready?





# Features and filters





# Features and filters

- Detecting distinctive + repeatable features
- Describing images with local statistics







# Indexing and search



- Matching features and regions across images


# Image formation

- How does light in 3d world project to form 2d images?




# Multiple views

- Multi-view geometry, matching, invariant features, stereo vision




a



Hartley and Zisserman




Lowe




Fei-Fei Li




# Grouping and fitting




Parallelism



Symmetry



Continuity






Closure







[fig from Shi et al]

- Clustering, segmentation, fitting; what parts belong together?








# Visual recognition



- Recognizing objects and categories, learning techniques

# Object detection

- Detecting novel instances of objects
- Classifying regions as one of several categories





# Attribute-based description

- Describing the high-level properties of objects
- Allows recognition of unseen objects



Naming

Aeroplane

otter

black: yes  
white: no  
brown: yes  
stripes: no  
water: yes  
eats fish: yes



Description

Unknown  
Has Wheel  
Has Wood

polar bear

black: no  
white: yes  
brown: no  
stripes: no  
water: yes  
eats fish: yes



Unusual attributes

Bird  
No Head  
No Beak

zebra

black: yes  
white: yes  
brown: no  
stripes: yes  
water: no  
eats fish: no




Unexpected attributes

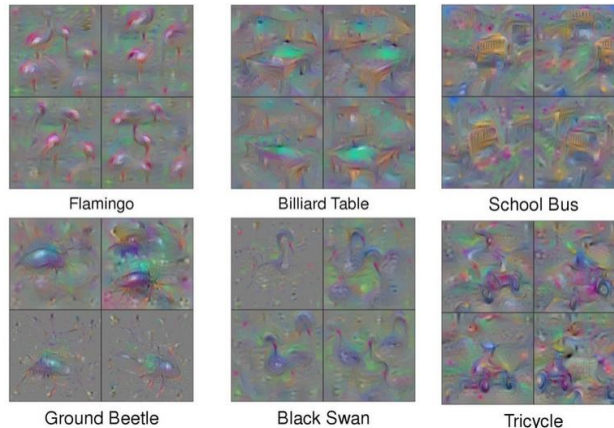
Motorbike  
Has Cloth

# Convolutional neural networks

- State-of-the-art on many recognition tasks



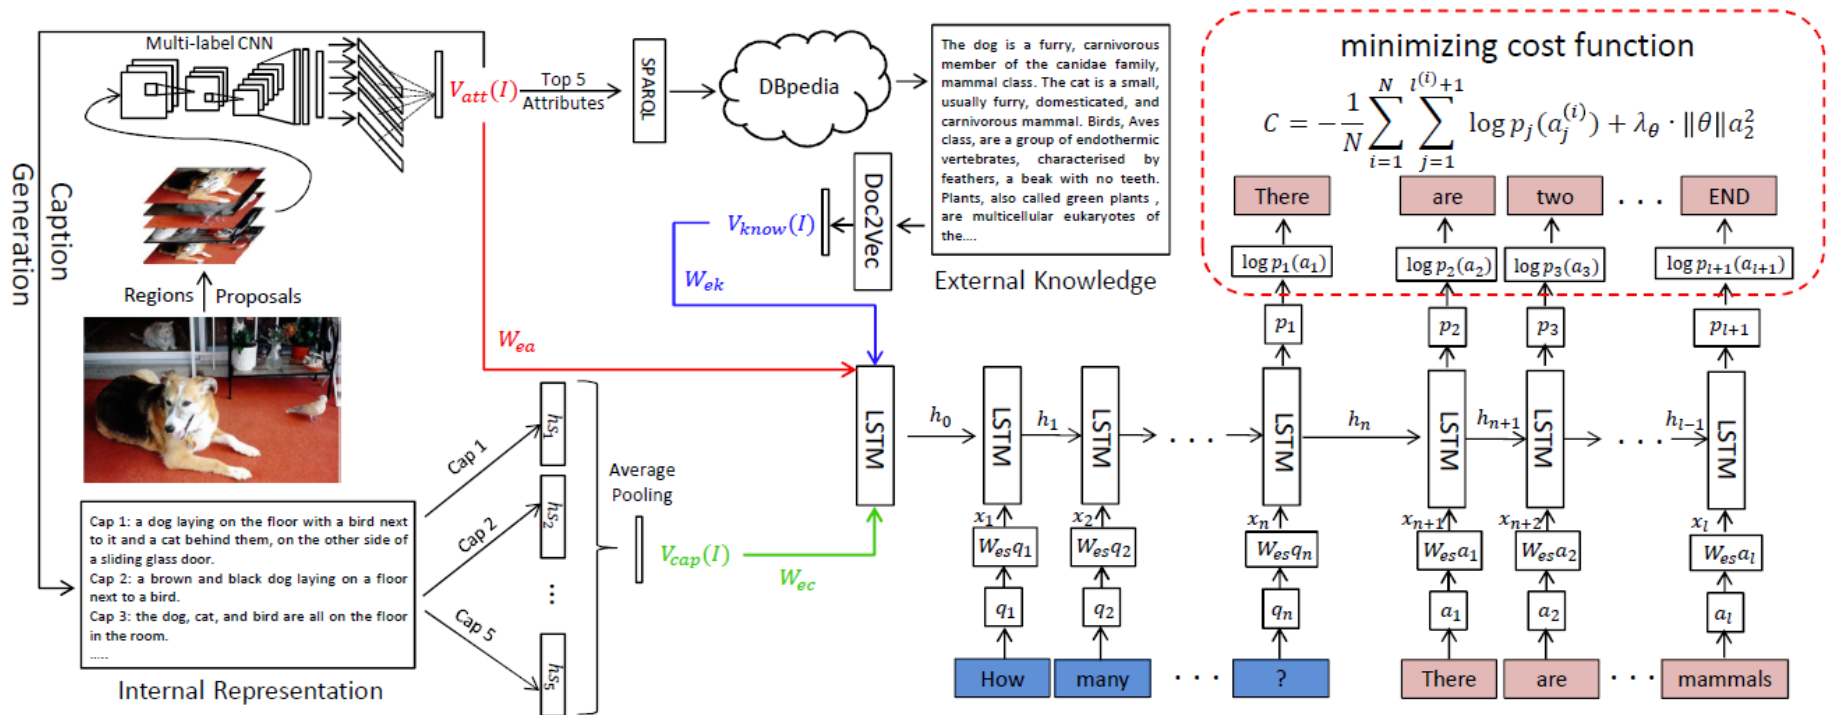
Krizhevsky et al.



Yosinski et al., ICML DL workshop 2015


# Recurrent neural networks

- Sequence processing, e.g. question answering



# Motion and tracking

- Tracking objects, video analysis




Tomas Izo



# Pose and actions

- Automatically annotating human pose (joints)
- Recognizing actions in first-person video




What are images?

# What are images? (in Matlab)

- Matlab treats images as matrices of numbers
- To proceed, let's talk very briefly about how images are formed



# Image formation



# Digital camera




A digital camera replaces film with a sensor array

- Each cell in the array is light-sensitive diode that converts photons to electrons

<http://electronics.howstuffworks.com/cameras-photography/digital/digital-camera.htm>

# Digital images



a b


**FIGURE 2.17** (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.



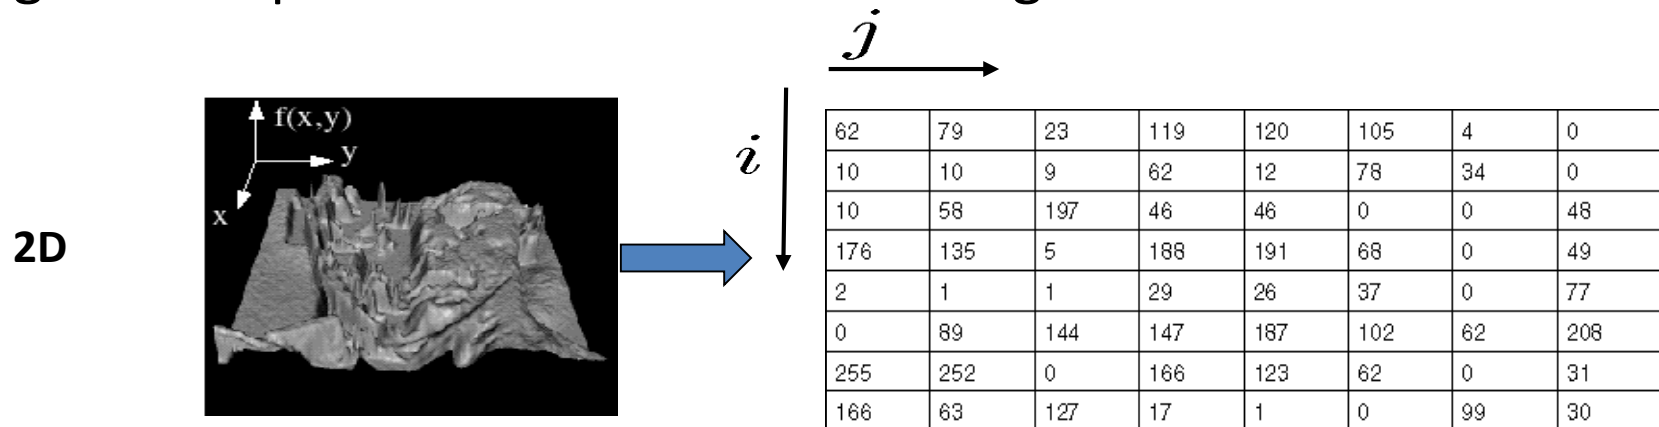
- **Sample** the 2D space on a regular grid
- **Quantize** each sample (round to nearest integer)

# Digital images


- **Sample** the 2D space on a regular grid
- **Quantize** each sample (round to nearest integer)
- What does quantizing signal look like?



- Image thus represented as a matrix of integer values.



# Digital color images



**Bayer filter**

© 2000 How Stuff Works

# Digital color images

Color images,  
RGB color space:

Split image into  
three channels



R



G



B





# Review and Tutorial

- Linear algebra
  - Very brief, all you need to know for most of course
  - Exception: Last three lectures before midterm (some more review then)
  - Raise your hand if you've had a linear algebra course
- Matlab:  
<http://www.cs.pitt.edu/~kovashka/cs1674/tutorial.m>

# Vectors and Matrices

- Vectors and matrices are just collections of ordered numbers that represent something: movements in space, scaling factors, word counts, movie ratings, pixel brightnesses, etc.
- We'll define some common uses and standard operations on them.

# Vector

- A column vector  $\mathbf{v} \in \mathbb{R}^{n \times 1}$  where

$$\mathbf{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix}$$

- A row vector  $\mathbf{v}^T \in \mathbb{R}^{1 \times n}$  where


$$\mathbf{v}^T = [v_1 \quad v_2 \quad \dots \quad v_n]$$

$T$  denotes the transpose operation

# Vector

- You'll want to keep track of the orientation of your vectors when programming in MATLAB.
- You can transpose a vector  $V$  in MATLAB by writing  $V'$ .

# Vectors have two main uses



- Vectors can represent an offset in 2D or 3D space
- Points are just vectors from the origin
- Data can also be treated as a vector
- Such vectors don't have a geometric interpretation, but calculations like "distance" still have value

# Matrix

- A matrix  $\mathbf{A} \in \mathbb{R}^{m \times n}$  is an array of numbers with size  $m \downarrow$  by  $n \rightarrow$ , i.e.  $m$  rows and  $n$  columns.

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} & \dots & a_{1n} \\ a_{21} & a_{22} & a_{23} & \dots & a_{2n} \\ \vdots & & & & \vdots \\ a_{m1} & a_{m2} & a_{m3} & \dots & a_{mn} \end{bmatrix}$$

- If  $m = n$ , we say that  $\mathbf{A}$  is square.

# Matrix Operations

- Addition

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} + \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} a+1 & b+2 \\ c+3 & d+4 \end{bmatrix}$$

- Can only add matrices with matching dimensions, or a scalar to a matrix.

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} + 7 = \begin{bmatrix} a+7 & b+7 \\ c+7 & d+7 \end{bmatrix}$$

- Scaling

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \times 3 = \begin{bmatrix} 3a & 3b \\ 3c & 3d \end{bmatrix}$$

# Matrix Operations


- Inner product (*dot* · product) of vectors
  - Multiply corresponding entries of two vectors and add up the result
  - We won't worry about the geometric interpretation for now

$$\mathbf{x}^T \mathbf{y} = \begin{bmatrix} x_1 & \dots & x_n \end{bmatrix} \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} = \sum_{i=1}^n x_i y_i \quad (\text{scalar})$$



# Matrix Operations

- Multiplication
- The product  $AB$  is:



- Each entry in the result is (that row of  $A$ ) dot product with (that column of  $B$ )

# Matrix Operations

- Multiplication example:

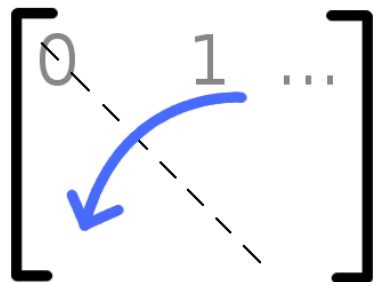
$$\begin{array}{ccc} A & \times & B \\ \downarrow & & \searrow \\ \begin{bmatrix} 0 & 2 \\ 4 & 6 \end{bmatrix} & & \begin{bmatrix} 1 & 3 \\ 5 & 7 \end{bmatrix} \\ & & \begin{bmatrix} \square & 14 \\ \square & \square \end{bmatrix} \end{array}$$

$$0 \cdot 3 + 2 \cdot 7 = 14$$

- Each entry of the matrix product is made by taking the dot product of the corresponding row in the left matrix, with the corresponding column in the right one.

# Matrix Operations

- Transpose – flip matrix, so row 1 becomes column 1


$$\begin{bmatrix} 0 & 1 & \dots \\ 2 & 3 & \\ 4 & 5 & \end{bmatrix}^T = \begin{bmatrix} 0 & 2 & 4 \\ 1 & 3 & 5 \end{bmatrix}$$

- A useful identity:

$$(ABC)^T = C^T B^T A^T$$

# Special Matrices

- Identity matrix  $\mathbf{I}$ 
  - Square matrix, 1's along diagonal, 0's elsewhere
  - $\mathbf{I} \cdot [\text{another matrix}] = [\text{that matrix}]$

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- Diagonal matrix
  - Square matrix with numbers along diagonal, 0's elsewhere
  - A diagonal  $\cdot$  [another matrix] scales the rows of that matrix

$$\begin{bmatrix} 3 & 0 & 0 \\ 0 & 7 & 0 \\ 0 & 0 & 2.5 \end{bmatrix}$$

1-Minute Break

# Matlab Tutorial

<http://www.cs.pitt.edu/~kovashka/cs1674/tutorial.m>

We'll cover parts 1-4, do parts 5-6 at home.

# Homework 1W and 1P

<http://people.cs.pitt.edu/~kovashka/cs1674/hw1w.html>

<http://people.cs.pitt.edu/~kovashka/cs1674/hw1p.html>

# Next Time / Homework

- Finish the Matlab tutorial on your own + post on Piazza if questions
- No class Monday (Labor Day), but HW1W due
- Wednesday: Image filtering, HW1P due
- Reading for Wednesday: Szeliski Sec. 3.2
- Reminder: Fill out Doodle!